

Advertising on the Web

互联网广告

Online Algorithms

- **Classic model of algorithms**


- You get to see the entire input, then compute some function of it
- In this context, “offline algorithm”

- **Online Algorithms**

- You get to see the input one piece at a time, and need to make irrevocable decisions along the way
- **Similar to the data stream model**

Online Bipartite Matching


Example: Bipartite Matching



Nodes: Boys and Girls; Edges: Preferences


Goal: Match boys to girls so that maximum number of preferences is satisfied

Example: Bipartite Matching



$M = \{(1,a), (2,b), (3,d)\}$ is a **matching**
Cardinality of matching = $|M| = 3$

Example: Bipartite Matching



$M = \{(1,c), (2,b), (3,d), (4,a)\}$ is a
perfect matching

Perfect matching ... all vertices of the graph are matched

Maximum matching ... a matching that contains the largest possible number of matches


Matching Algorithm

- **Problem:** Find a maximum matching for a given bipartite graph
 - A perfect one if it exists
- There is a polynomial-time offline algorithm based on augmenting paths (Hopcroft & Karp 1973, see http://en.wikipedia.org/wiki/Hopcroft-Karp_algorithm)
- **But what if we do not know the entire graph upfront?**

Online Graph Matching Problem

- Initially, we are given the set boys
- In each round, one girl's choices are revealed
 - That is, girl's edges are revealed
- At that time, we have to decide to either:
 - Pair the girl with a boy
 - Do not pair the girl with any boy
- Example of application:
Assigning tasks to servers

Online Graph Matching: Example



(1,a)
(2,b)
(3,d)

Greedy Algorithm

- Greedy algorithm for the online graph matching problem:
 - Pair the new girl with **any** eligible boy
 - If there is none, do not pair girl
- How good is the algorithm?

Competitive Ratio

- For input I , suppose greedy produces matching M_{greedy} while an optimal matching is M_{opt}

Competitive ratio =


$$\min_{\text{all possible inputs } I} (|M_{greedy}| / |M_{opt}|)$$

(what is greedy's worst performance over all possible inputs I)

Analyzing the Greedy Algorithm

- Consider a case: $M_{greedy} \neq M_{opt}$
- Consider the set G of girls matched in M_{opt} but not in M_{greedy}
- Then every boy B adjacent to girls in G is already matched in M_{greedy} :
 - If there would exist such non-matched (by M_{greedy}) boy adjacent to a non-matched girl then greedy would have matched them
- Since boys B are already matched in M_{greedy} then


(1) $|M_{greedy}| \geq |B|$




Analyzing the Greedy Algorithm

■ Summary so far:

- Girls G matched in M_{opt} but not in M_{greedy}
 - (1) $|M_{greedy}| \geq |B|$
- There are at least $|G|$ such boys ($|G| \leq |B|$) otherwise the optimal algorithm couldn't have matched all girls in G
 - So: $|G| \leq |B| \leq |M_{greedy}|$
- By definition of G also: $|M_{opt}| \leq |M_{greedy}| + |G|$
 - Worst case is when $|G| = |B| = |M_{greedy}|$
- $|M_{opt}| \leq 2|M_{greedy}|$ then $|M_{greedy}|/|M_{opt}| \geq 1/2$



Worst-case Scenario



(1,a)
(2,b)

Web Advertising

History of Web Advertising

- **Banner ads (1995-2001)**

- Initial form of web advertising

- Popular websites charged X\$ for every 1,000 “impressions” of the ad

- Called “**CPM**” rate
(Cost per thousand impressions)

- Modeled similar to TV, magazine ads

- From **untargeted** to **demographically targeted**

- **Low click-through rates**

- Low ROI for advertisers



CPM...cost per mille
Mille...thousand in Latin

Performance-based Advertising

- Introduced by Overture around 2000
 - Advertisers **bid on search keywords**
 - When someone searches for that keyword, the **highest bidder's ad is shown**
 - Advertiser is charged only if the ad is clicked on
- Similar model adopted by Google with some changes around 2002
 - Called **Adwords**

Ads vs. Search Results

Web

Results 1 - 10 of about 2,230,000 for **geico**. (0.04 sec)

[GEICO](#) Car Insurance. Get an auto insurance quote and save today ...

GEICO auto insurance, online car insurance quote, motorcycle insurance quote, online insurance sales and service from a leading insurance company.

[www.geico.com/](#) - 21k - Sep 22, 2005 - [Cached](#) - [Similar pages](#)

[Auto Insurance](#) - [Buy Auto Insurance](#)

[Contact Us](#) - [Make a Payment](#)

[More results from www.geico.com »](#)

[Geico](#), Google Settle Trademark Dispute

The case was resolved out of court, so advertisers are still left without legal guidance on use of trademarks within ads or as keywords.

[www.clickz.com/news/article.php/3547356](#) - 44k - [Cached](#) - [Similar pages](#)

Google and [GEICO](#) settle AdWords dispute | The Register

Google and car insurance firm **GEICO** have settled a trade mark dispute over ... Car insurance firm **GEICO** sued both Google and Yahoo! subsidiary Overture in ...

[www.theregister.co.uk/2005/09/09/google_geico_settlement/](#) - 21k - [Cached](#) - [Similar pages](#)

[GEICO](#) v. Google

... involving a lawsuit filed by Government Employees Insurance Company (**GEICO**). **GEICO** has filed suit against two major Internet search engine operators, ...

[www.consumeraffairs.com/news04/geico_google.html](#) - 19k - [Cached](#) - [Similar pages](#)

Sponsored Links

[Great Car Insurance Rates](#)

Simplify Buying Insurance at Safeco
See Your Rate with an Instant Quote
[www.Safeco.com](#)

[Free Insurance Quotes](#)

Fill out one simple form to get multiple quotes from local agents.
[www.HometownQuotes.com](#)

[5 Free Quotes. 1 Form.](#)

Get 5 Free Quotes In Minutes!
You Have Nothing To Lose. It's Free
[sayyessoftware.com/Insurance](#)
Missouri

Web 2.0

- **Performance-based advertising works!**
 - Multi-billion-dollar industry
- **Interesting problem:**
What ads to show for a given query?
 - (Today's lecture)
- **If I am an advertiser, which search terms should I bid on and how much should I bid?**
 - (Not focus of today's lecture)

Adwords Problem

■ Given:

- 1. A set of bids by advertisers for search queries
- 2. A click-through rate for each advertiser-query pair
- 3. A budget for each advertiser (say for 1 month)
- 4. A limit on the number of ads to be displayed with each search query

■ Respond to each search query with a set of advertisers such that:

- 1. The size of the set is no larger than the limit on the number of ads per query
- 2. Each advertiser has bid on the search query
- 3. Each advertiser has enough budget left to pay for the ad if it is clicked upon

Adwords Problem

- A stream of queries arrives at the search engine: q_1, q_2, \dots
- Several advertisers bid on each query
- When query q_i arrives, search engine must pick a subset of advertisers whose ads are shown
- **Goal:** Maximize search engine's revenues
 - **Simple solution:** Instead of raw bids, use the “expected revenue per click” (i.e., $\text{Bid} \times \text{CTR}$)
- **Clearly we need an online algorithm!**

The Adwords Innovation

Advertiser	Bid	CTR	Bid * CTR
A	\$1.00	1%	1 cent
B	\$0.75	2%	1.5 cents
C	\$0.50	2.5%	1.125 cents

Click through
rate

Expected
revenue

The Adwords Innovation

Advertiser	Bid	CTR	Bid * CTR
B	\$0.75	2%	1.5 cents
C	\$0.50	2.5%	1.125 cents
A	\$1.00	1%	1 cent

Complications: Budget

- **Two complications:**
 - Budget
 - CTR of an ad is unknown
- **Each advertiser has a limited budget**
 - Search engine guarantees that the advertiser will not be charged more than their daily budget

Complications: CTR

- **CTR: Each ad has a different likelihood of being clicked**
 - **Advertiser 1** bids \$2, click probability = 0.1
 - **Advertiser 2** bids \$1, click probability = 0.5
 - **Clickthrough rate (CTR)** is measured **historically**
 - **Very hard problem: Exploration vs. exploitation**
 - Exploit:** Should we keep showing an ad for which we have good estimates of click-through rate
 - or**
 - Explore:** Shall we show a brand new ad to get a better sense of its click-through rate

Greedy Algorithm

- **Our setting: Simplified environment**
 - There is **1** ad shown for each query
 - All advertisers have the same budget **B**
 - All ads are equally likely to be clicked
 - Value of each ad is the same (**$=1$**)
- **Simplest algorithm is greedy:**
 - For a query pick any advertiser who has bid **1** for that query
 - **Competitive ratio of greedy is $1/2$**

Bad Scenario for Greedy

- **Two advertisers A and B**
 - A bids on query x , B bids on x and y
 - Both have budgets of \$4
- **Query stream: $x x x x y y y y$**
 - Worst case greedy choice: $B B B B _ _ _ _$
 - Optimal: $A A A A B B B B$
 - **Competitive ratio = $\frac{1}{2}$**
- **This is the worst case!**
 - **Note:** Greedy algorithm is deterministic – it always resolves draws in the same way

BALANCE Algorithm [MSVV]

- **BALANCE** Algorithm by Mehta, Saberi, Vazirani, and Vazirani
 - For each query, pick the advertiser with the largest unspent budget
 - Break ties arbitrarily (but in a deterministic way)


Example: BALANCE

- **Two advertisers A and B**
 - A bids on query x , B bids on x and y
 - Both have budgets of \$4
- **Query stream:** $x x x x y y y y$
- **BALANCE choice:** A B A B B B _ _
 - Optimal: A A A A B B B B
- **In general:** For **BALANCE** on 2 advertisers
Competitive ratio = $\frac{3}{4}$

Analyzing BALANCE

- **Consider simple case (w.l.o.g.):**
 - 2 advertisers, A_1 and A_2 , each with budget B (≥ 1)
 - Optimal solution exhausts both advertisers' budgets
- **BALANCE must exhaust at least one advertiser's budget:**
 - **If not, we can allocate more queries**
 - Whenever BALANCE makes a mistake (both advertisers bid on the query), advertiser's unspent budget only decreases
 - Since optimal exhausts both budgets, one will for sure get exhausted
 - Assume BALANCE exhausts A_2 's budget, but allocates x queries fewer than the optimal
 - **Revenue: $BAL = 2B - x$**

Analyzing Balance




■ Queries allocated to A_1 in the optimal solution

■ Queries allocated to A_2 in the optimal solution

Optimal revenue = $2B$

Assume Balance gives revenue = $2B - x = B + y$



Unassigned queries should be assigned to A_2
(if we could assign to A_1 we would since we still have the budget)

Goal: Show we have $y \geq x$

Case 1) $\leq \frac{1}{2}$ of A_1 's queries got assigned to A_2
then $y \geq B/2$


Case 2) $> \frac{1}{2}$ of A_1 's queries got assigned to A_2
then $x \leq B/2$ and $x + y = B$

Balance revenue is minimum for $x = y = B/2$

Minimum Balance revenue = $3B/2$

Competitive Ratio = $3/4$

BALANCE exhausts A_2 's budget



BALANCE: General Result

- In the general case, worst competitive ratio of BALANCE is $1 - 1/e = \text{approx. } 0.63$
 - Interestingly, no online algorithm has a better competitive ratio!
- Let's see the worst case example that gives this ratio


Worst case for BALANCE

- **N advertisers:** A_1, A_2, \dots, A_N
 - Each with budget $B > N$
- **Queries:**
 - $N \cdot B$ queries appear in N rounds of B queries each
- **Bidding:**
 - Round 1 queries: bidders A_1, A_2, \dots, A_N
 - Round 2 queries: bidders A_2, A_3, \dots, A_N
 - Round i queries: bidders A_i, \dots, A_N
- **Optimum allocation:**

Allocate round i queries to A_i

 - Optimum revenue $N \cdot B$

BALANCE Allocation

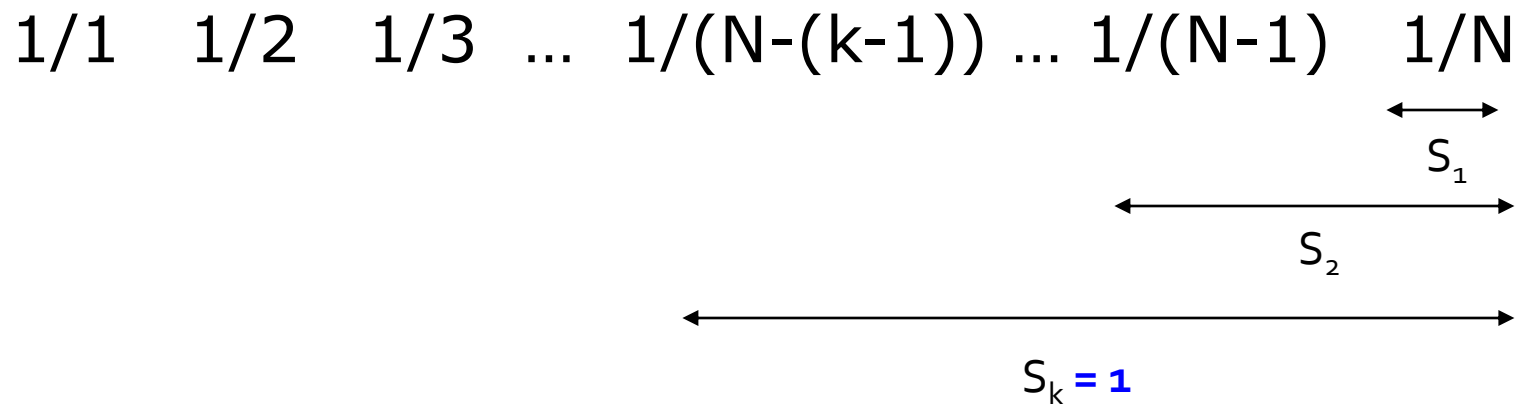



BALANCE assigns each of the queries in round 1 to N advertisers. After k rounds, sum of allocations to each of advertisers A_k, \dots, A_N is

$$S_k = S_{k+1} = \dots = S_N = \sum_{i=1}^{k-1} \frac{B}{N-(i-1)}$$


If we find the smallest k such that $S_k \geq B$, then after k rounds we cannot allocate any queries to any advertiser

BALANCE: Analysis



BALANCE: Analysis

- **Fact:** $H_n = \sum_{i=1}^n 1/i \approx \ln(n)$ for large n
 - Result due to Euler



- $S_k = 1$ implies: $H_{N-k} = \ln(N) - 1 = \ln\left(\frac{N}{e}\right)$
 - We also know: $H_{N-k} = \ln(N-k)$
 - So: $N - k = \frac{N}{e}$
 - Then: $k = N\left(1 - \frac{1}{e}\right)$
- N terms sum to $\ln(N)$.
Last k terms sum to 1.
First $N-k$ terms sum to $\ln(N-k)$ but also to $\ln(N)-1$*

BALANCE: Analysis

- So after the first $k=N(1-1/e)$ rounds, we cannot allocate a query to any advertiser
- Revenue = $B \cdot N (1-1/e)$
- Competitive ratio = $1-1/e$

General Version of the Problem

- **Arbitrary bids and arbitrary budgets!**
- Consider we have 1 query q , advertiser i
 - Bid = x_i
 - Budget = b_i
- **In a general setting BALANCE can be terrible**
 - Consider two advertisers A_1 and A_2
 - $A_1: x_1 = 1, b_1 = 110$
 - $A_2: x_2 = 10, b_2 = 100$
 - Consider we see **10** instances of q
 - BALANCE always selects A_1 and earns **10**
 - Optimal earns **100**

Generalized BALANCE

- **Arbitrary bids:** consider query q , bidder i
 - Bid = x_i
 - Budget = b_i
 - Amount spent so far = m_i
 - Fraction of budget left over $f_i = 1 - m_i/b_i$
 - Define $\psi_i(q) = x_i(1 - e^{-f_i})$
- Allocate query q to bidder i with largest value of $\psi_i(q)$
- **Same competitive ratio $(1 - 1/e)$**